

Harnessing cloud computing for high capacity analysis of neuroimaging data

Daniel Clark¹, Christian Haselgrove³, David N. Kennedy³, Zhizhong Liu⁴, Michael Milham^{1,5}, Petros Petrosyan⁴, Carinna M. Torgerson⁴, John D. Van Horn⁴, and R. Cameron Craddock^{1,2,*}

¹Center for the Developing Brain, Child Mind Institute, New York, NY, USA

²Computational Neuroimaging Laboratory, Center for Biomedical Imaging and Neuromodulation, Nathan S. Kline Institute for Psychiatric Research, Orangeburg, NY, USA

³Division of Informatics, Department of Psychiatry, University of Massachusetts Medical School, Worcester, MA, USA

⁴The Institute for Neuroimaging and Informatics (INI) and Laboratory of Neuro Imaging (LONI), Keck School of Medicine of USC, University of Southern California, Los Angeles, CA, USA

⁵Center for Biomedical Imaging and Neuromodulation, Nathan S. Kline Institute for Psychiatric Research, Orangeburg, New York, USA

Correspondence*:

R. Cameron Craddock

Computational Neuroimaging Laboratory, Center for Biomedical Imaging and Neuromodulation, Nathan S. Kline Institute for Psychiatric Research, Orangeburg, NY, 10962, USA, email@uni.edu

ABSTRACT

For full guidelines regarding your manuscript please refer to Author Guidelines or **Table 1** for a summary according to article type.

Keywords: Cloud computing, high-performance compute clusters, neuroimaging preprocessing, AWS, spot instances

1 INTRODUCTION

intro

2 METHODS

2.1 The Amazon Web Services Elastic Compute Cloud

Amazon provides a collection of remote computing services called Amazon Web Services (AWS). One of the most fundamental and well-known of these services is Amazon Elastic Compute Cloud (EC2). EC2 allows for the configuration and use of virtual servers, or compute instances, at any capacity desired. In contrast with the up-front research and costs necessary to purchase an in-house computing solution, EC2 offers the flexibility to utilize a wide range of computing capacity with a "pay-as-you-go" model. The associated costs are divided between processing/compute hours, storage amount, and data

transfer/networking traffic. AWS is offered across multiple regions hosted at different data centers spread out over the world with each region subdivided into availability zones. For example the us-east-1 region is centered in northern Virginia and its services are subdivided among the us-east-1a, us-east-1b, us-east-1c, and us-east-1e availability zones. This framework makes it possible to optimize one's computing power, budget, and locality for their application.

2.1.1 Instances and AMIs

EC2 offers “instances” which represent the hardware resources that will perform computation. To deliver consistent performance, the hardware underlying instances are bundled by Amazon into virtual processors, called Amazon EC2 Compute Units (ECUs). ECUs are routinely benchmarked and tested to ensure stability even as hardware is modified or updated. A range of “instance types” are available to support a variety of different use cases. Instance types include “general purpose” instances (e.g. t2.micro, t2.small, m4.xlarge), with low to moderate computing power, “compute optimized” instances (e.g. c3.large, c3.8xlarge, c4.4xlarge) that have a higher quantity of CPUs, “memory optimized” instances (e.g. m3.2xlarge, r3.xlarge, r3.8xlarge) containing large amounts of RAM, “storage optimized” instances (e.g. i2.xlarge, i2.4xlarge, d2.8xlarge) for high data I/O throughput and capacity, and GPU instances (e.g. g2.2xlarge, g2.8xlarge). As of June 2015, the “compute optimized” instances run with Intel Xeon E5-2666 v3 (Haswell) processors.

The software ecosystem to be run on an instance is provided by an Amazon Machine Image (AMI). AMIs are analogous to virtual machine images and contain the operating system, software, and data for a virtual server. Users can choose from thousands of preconfigured AMIs from the AWS and community marketplaces or can create their own. AMIs are available using Linux or Windows operating systems and are priced for hourly or yearly usage; many Linux-based AMIs are free.

Preconfigured AMIs provide users with a platform-as-a-service in which they rely on the developer to configure and maintain the software ecosystem. This allows developers to focus on maintaining one central system that can be deployed as needed, rather than keeping multiple resources up-to-date. This solution is novel as it allows for centralized system administration with dynamically-sized “elastic” computing resources; end-users only pay for what they use and without having to perform configuration or maintenance themselves.

2.1.1.1 On-demand instances

One of the ways to pay for and use EC2 instances is via on-demand pricing. In this schema, the user can request their desired resources at any time by paying the on-demand hourly rates. The instances are available to them immediately without interruptions and will continue to be active until explicitly terminated. The cost of on-demand instances is significantly higher than other purchasing mediums, but is useful for time-sensitive tasks.

2.1.1.2 Reserved instances

Another payment system is to use reserved instances. If there is long term interest in using EC2 and demand for resources is pretty consistent over time, users can pay up-front, or throughout the reservation duration (1 year or 3 years) to reserve instances to be used at any time. The cost is reduced resulting in savings of up to 75% compared with on-demand instances.

Table 1. Instance types pricing

Instance type	vCPUs	ECUs	Memory (GiB)	On-demand (\$/hr)	Reserved no-upfront, 1 year (\$/hr)	Typical spot (\$/hr)
t2.small	1	Variable	2	0.026	0.018	N/A
m3.xlarge	4	13	15	0.266	0.190	0.0334
c3.8xlarge	32	108	60	1.680	1.168	0.350
g2.2xlarge	8	26	15	0.650	0.474	0.075

*All prices reflect the us-east-1 northern Virginia region

2.1.1.3 Spot instances

Amazon allows for users to bid on the resources that are not being used at a significant discount, up to 90%, compared to on-demand pricing. The instances themselves operate in the exact same way as an on-demand or reserved instance would; the major difference with spot is the chance of forced termination if the spot price equals or exceeds the user bid price. With spot pricing, the hourly cost for an instance can vary depending on market demand. This provides great flexibility for cost and computing demands as long as the type of processing can handle sudden interruptions and recover.

2.1.2 Data storage

2.1.2.1 Elastic block storage

Data storage pricing is offered for two storage mediums in EC2: elastic block storage (EBS) and instance storage. EBS offers block-level storage volumes that can be treated like disk drives that support formatting of file systems (such as NFS or NTFS). These volumes can be attached to running instances and persist even after the instance is shut down or terminated. EBS volumes are billed based on GB-months of storage, IOPS, and I/O requests, depending on the type of EBS storage used.

2.1.2.2 Instance storage

Instance storage is also a block-level storage medium available in EC2, however, it is only available with certain instance types, at certain sizes, and only persists as long as the instance is up and running. Instance storage is located on disks that are physically attached to the host instance and will lose all data when the instance is stopped or terminated. The pricing for instance storage is built into the hourly price of using the associated instance - this makes it a cheaper option than EBS if data need only persist as long as the instances. Instance storage also provides for optimizing storage costs by keeping unnecessary or temporary files with the instance, and storing valuable files on smaller EBS volumes.

2.1.2.3 S3

Amazon also offers their Simple Storage Service (S3) platform as a cost efficient and robust alternative to solutions in EC2. S3 uses object-level storage in user-defined buckets, rather than block-level storage and can be accessed from anywhere via their web interface or API. S3 storage comes in standard, reduced redundancy, and glacier storage costs - depending on the importance or access time needed of the data. Standard storage offers the most readily accessible and robust data with 99.999999999% durability, where reduced redundancy offers slightly less robustness with 99.99% durability, but at a reduced cost. Glacier is Amazons data archiving service that allows users to store backups of their S3 data at a further reduced cost, with the tradeoff of waiting 3 to 5 hours to access the data again.

All storage services prices are based on per-GB of data at various tiers of overall storage being utilized in S3 (e.g. \$0.03 per GB for up to 1TB of overall storage per month); these prices vary between regions.

Additional costs are associated with the amount of requests (e.g. PUT, DELETE, COPY, LIST) made on the S3 objects. However, the costs are much lower compared with long-term data storage in EC2.

2.1.3 Data transfer

There are costs associated with the amount of data transferred in and out of AWS. Data transfer in to EC2 from anywhere on the internet is completely free for all regions. However, there are costs associated with transferring data out to the internet (priced per GB), depending on how much data in total is transferred monthly. Data transfer is free between instances in EC2 as long as they are in the same availability zone and use private IP addresses.

Data transfer pricing has a fairly simple structure in S3. It is free to upload data to S3 from anywhere. Costs are associated for every GB of data downloaded from S3 to the internet in a similar monthly total tier that EC2 uses. However, downloading data from S3 is free to EC2 instances in the same AWS region as the S3 bucket of interest.

2.1.4 Databases

Amazon offers a variety of database solutions in AWS, each of which are tailored for specific use cases. For relational databases, Amazon Relational Database Service (RDS) provides a range of database engines, including MySQL, Oracle, PostgreSQL, and more. RDS abstracts away the administrative overhead with setting up a database and provides automatic backups and elastic compute resources using their API.

Other solutions include DynamoDB, for NoSQL databases with a lower and more predictable workload, ElastiCache, for in-memory cache rapid access of data, and Redshift, for high-performance big-data warehousing. Additionally, users can host their choice of database engine, with any configuration, on EC2 instances.

2.1.5 Data privacy, security, and HIPAA

Amazon has built a robust security infrastructure around their web services, with many features being specifically targeted for regulatory compliance. Security administration in AWS is mostly the same as it is with in-house servers, with security patches, backups, anti-virus software, user permissions, ACLs, traffic monitoring, and data encryption. Hardware VPNs can be set up between local and cloud resources as well. There are a few administrative differences with regard to security, as the resources are being managed by the user remotely. These include utilizing software-based security instead of hardware-based approaches, issuing user credential keys for access, configuring private networks using AWS Virtual Private Cloud (VPC), setting up firewalls around each EC2 instance in the form of security groups, and dealing with geographical isolation between regions. It is important to note that security is now a shared responsibility between the user and AWS. The security associated with hardware, networking, facilities, and infrastructure software is managed by Amazon. The robustness and security of these are top priority and are routinely audited by third parties for compliance. The user is responsible for who has access to their resources, software security patches, data encryption, network firewalls, and backup and archiving preferences; overall the user is responsible for the proper configuration of the AWS infrastructure to keep data secure and private. This relationship is outlined in detail in the “AWS Shared Responsibility Model.”

Utilizing AWS resources to process and store PHI data requires HIPAA and HITECH compliance; as such, any institutions dealing with PHI needs to be HIPAA-certified. HIPAA certification is not specifically available for cloud computing providers like AWS. However, AWS has aligned their risk management program with FedRAMP and NIST 800-53 security standards; NIST (the National Institute of Standards

and Technology) has supported this as a valid approach for satisfying the HIPAA security rule, which states “the EPHI that a covered entity creates, receives, maintains, or transmits must be protected against reasonably anticipated threats, hazards, and impermissible uses and/or disclosures.” This allows for any HIPAA-certified institution to use AWS services for the processing, transmission, and storage of PHI data, as long as they take the proper precautions.

HIPAA rules require covered institutions enter into a contract with any non-covered subcontractors or businesses to guarantee both are aware and take precautions to guard the integrity of PHI; this contract is known as a “Business Associate Agreement” (BAA). Amazon offers a standard BAA and will sign with HIPAA-covered institutions, given that both parties obey the AWS Shared Responsibility Model.

Once the user has signed a BAA with Amazon, all AWS services are available to them, but only six services should be used when dealing with PHI: Amazon EC2, Amazon EBS, Amazon S3, Amazon Redshift, Amazon Glacier, and Amazon Elastic Load Balancer. AWS has set up each of these services to provide for any and all HIPAA safeguards in their configurations; the default configurations on any of these services are typically the most stringent and secure. The responsibilities fall to the user to ensure proper user access and permissions, which can be controlled with AWS Identity and Access Management (IAM) service, firewalls, and traffic restrictions to specific IP ranges. System administrators interacting with AWS can grant users access with key or token credentials. Additionally, PHI data should always be encrypted before, during, and after transfer to AWS in accordance with HIPAA regulations; this can be achieved using technologies such as 256-bit AES encryption algorithms. Finally, HIPAA requires robust data-backup and disaster-recovery procedures as well as the ability to audit systems for their security and PHI data provenance. AWS offers the ability to chronologically snapshot any EBS volumes in EC2 and back up all relevant data to S3, where it is copied across data centers and securely stored. Additionally, AWS employees are not able to log into customer instances in EC2 and S3 data access is highly restricted. In the end, the easiest way to ensure HIPAA-compliance is to always use anonymized data and avoid the upload of PHI information.

Currently other HIPAA-covered entities are using the AWS infrastructure for their health care services and many HIPAA compliance case studies can be found on the AWS website.

2.2 Neuroimaging in the cloud

2.2.1 AMIs and software tools

AWS allows researchers in neuroimaging to process and analyze their data using the aforementioned computational and storage resources. Several organizations and laboratories have developed AMIs in EC2 as a platform-as-a-service to offer neuroimaging pipelines to the general public. These include the Configurable-Pipeline for the Analysis of Connectomes (C-PAC), the Neuroimaging Informatics Tools and Resources Clearinghouse (NITRC) Computational Environment (CE), Human Connectome Project (HCP), and the Laboratory of Neuro Imaging (LONI) AMIs. All of these are available either on the AWS or community marketplaces in EC2 for free use. Each AMI comes with the necessary software dependencies and system configuration to make neuroimaging analysis ready out-of-box and straightforward. The Debian-based AMIs use NeuroDebian as a software repository resource for neuroscience software; this makes installing and updating many common packages, like AFNI and FSL, easy for anyone.

Part of the advantage of using EC2 as a computing resource is the ability to dynamically scale the amount of resources needed. In particular, image preprocessing can leverage a cluster computing configuration as each image can be analyzed independently; this is also known as an embarrassingly parallel workload. It

becomes necessary for the user to easily launch an array of compute nodes that can process their imaging data in parallel and write their outputs to a common storage place. Starcluster is a tool designed for just that purpose. It abstracts away a lot of the clutter and detail that comes with configuring a cluster on EC2, and allows the user to easily configure and customize cluster templates to their need. By leveraging Starcluster with the CPAC, NITRC, and LONI AMIs, an interested user can accomplish high data throughput with minimal effort, time, and cost.

2.2.2 Data available in the cloud

In addition to the elastic computational resources cloud computing has to offer, it also presents a useful medium for sharing data. Currently, the 1000 Functional Connectomes Project and International Neuroimaging Data-sharing Initiative (FCP/INDI) host more than 11 TB of free, publicly available data for download and research use on S3. HCP has over 500 participants-worth of free, downloadable data.

The National Institutes of Health have been hosting data since the launch of the National Database for Autism Research (NDAR) project in 2006. NDAR encompasses neuroimaging and behavioral data of over 90,000 participants focused on the origins and effects of autism spectrum disorder (ASD). NDAR uses Amazons RDS and S3 services to give researchers the ability to query, analyze, and contribute heterogenous data in a common repository.

2.3 Cloud-based cluster configurations

The options for instance pricing on AWS EC2 provides the user with many choices in how to configure a compute cluster in the cloud. These configurations can be optimized for specific time and costs demands. In cases where time is pressing and costs are less of an issue, the user can launch a master node and a number of slave nodes, all as on-demand, or reserved instances. This guarantees that all nodes of the cluster are up-and-running, promptly, and without worrying about interruptions or data loss. The trade-off to this approach is a higher cost.

Alternatively, in cases that are more cost-sensitive, users can populate cluster nodes as spot instances. By setting an appropriate bid price, the total cost can be mitigated while still achieving computing runs in a reasonable amount of time. The conventional approach here is to launch an on-demand, small to moderate sized master node, and larger spot-price compute nodes. This configuration provides some level of stability and continuity in the cases of node termination due to market demand as the master node can keep track of job status and event logging.

A factor that comes into either on-demand or spot cluster configurations is data storage. EBS volumes can be used as shared drives across the entire cluster; this provides for a way to incrementally store data as jobs are completing, from any node. The EBS volumes are also guaranteed to persist despite any changes in market prices or node terminations. However, provisioning a large-enough EBS volume to store an entire clusters worth of output data can be costly. Alternatively, intermediate data can be temporarily stored on each nodes instance storage. Many instance types come with a certain amount of GB built-in to the hourly price for use while the instance is up and running. Designing workflows to utilize this space to the fullest extent will minimize the necessary amount of persistent storage (whether its EBS or S3), and thus cost, needed.

The cluster configuration used for the benchmarks was taking advantage of the instance storage of each node for all intermediate data and saving outputs of interest to S3.

2.4 Benchmarks

Several institutions processed neuroimaging datasets, of various sizes, through different pipelines on EC2. The datasets include: the Autism Brain Imaging Data Exchange (ABIDE), which includes anatomical and resting state functional MRI (rfMRI) scans from 539 participants with autism spectrum disorders (ASD) and 573 typical controls, the Consortium for Reliability and Reproducibility (CoRR) dataset, which includes anatomical and rfMRI scans from 1,629 typical individuals, across 18 sites around the globe totaling in 5,093 scans, and 2,085 T1-weighted anatomical scans from NDAR.

All of the preprocessing was done for the NDAR data, including the ANTs runs, on an early revision of the main C-PAC AMI. This first AMI (ami-cc74e1a4) ran on a 64-bit Ubuntu 14.04 operating system and incorporated C-PAC v0.3.8 and all of its Python package dependencies, including the Python standard library, numpy, scipy, matplotlib, networkx, nipy, nibabel, traits, lockfile, yaml, jinja2, nose, pygraphviz, cython, ipython, and wxPython. Neuroimaging packages FSL, AFNI, and ANTs were also included. Additional features like the boto Python package, Oracle and the cx_Oracle Python package were installed as well; this made interaction with the miNDAR database and S3 buckets possible directly from the EC2 instance. Originally, this AMI was only available on the community marketplace in the us-east-1 region; the latest C-PAC AMI is available on the AWS marketplace across all regions and has the most up-to-date of these tools and is running C-PAC v0.3.9.1. All subsequent processing was done on this AMI.

Four pipelines were used in processing the aforementioned datasets: C-PAC, ANTs cortical thickness, Freesurfer recon-all (with and without GPU-enabled hardware), and the Quality Assessment Protocol (QAP). The C-PAC pipeline was run on the ABIDE dataset and the IBA_TRT site (50 subjects) from the CoRR dataset. This pipeline involved the structural and functional preprocessing using four noise-removal strategies: global signal removal with and without a 0.01 to 0.1 Hz band-pass filter, and non global signal removal with and without band-pass filtering. These images were registered to the MNI152 standard template and measured to produce a set of 19 statistical derivatives: ALFF, fALFF, REHO, 10 dual-regressed intrinsic connectivity networks, binarized and weighted degree centrality, binarized and weighted eigenvector centrality, lFCD, and VMHC. Additionally, time series were extracted and averaged over seven parcellated brain atlases, each of which containing multiple regions of interest (ROIs). Finally, the mean functional data, registered and noise-filtered functional data, and functional mask were saved as well.

The ANTs cortical thickness pipeline is a volume-based method for extracting cortical thickness estimates from anatomical MRI data. The pipeline is detailed here: <http://www.ncbi.nlm.nih.gov/pubmed/24879923>, and consists of nonlinear registration to a template, bias correction, tissue segmentation, and cortical thickness estimation. The mean cortical thickness was then calculated at 31 ROIs on each hemisphere of the cortex and using the Desikan-Killiany-Tourville (DKT-31) cortical labelling protocol. Afterwards, the normalized cortical thickness image, and mean ROI values were uploaded to the results bucket on S3. This analysis was done on the ABIDE dataset, IBA_TRT site of CoRR, and the 2,085 subjects analyzed from NDAR.

Freesurfer's recon-all pipeline extracts cortical and subcortical structures from anatomical data and computes surface and volumetric-based statistics. The pipeline is detailed here: <https://surfer.nmr.mgh.harvard.edu/fswiki/recon-all>. The image is motion and intensity-corrected, skull-stripped, normalized, registered, and segmented to produce many outputs, of which, statistical derivatives are calculated on. The recon-all function takes in an -openmp flag where the user can specify the number of cores to allocate to hyper-threaded executables underlying the pipeline; this decreases computation time.

The CoRR IBA_TRT site was ran through the recon-all pipeline, using 8 cores passed to the -openmp flag. Additionally, recon-all supports GPU hardware via the -use-cuda flag. By launching the g2.2xlarge instance type into a cluster configuration, recon-all was able to utilize CUDA-enabled executables on GPU hardware in the pipeline for further speed improvements; this pipeline also passed 8 cores to the -openmp flag for any executables that still needed to run on the CPU.

There are many proposed methods in scientific literature for analyzing the quality of MRI images. The quantitative methods have been assembled to form the Quality Assessment Protocol (QAP) pipeline. The QAP includes measures for gauging the quality of structural and functional MRI data. The structural measures include (taken from an excerpt of QAP Steve sent me): contrast to noise ratio (CNR; Magnotta and Friedman, 2006), entropy focus criterion (EFC, Atkinson 1997), foreground to background energy ratio (FBER), voxel smoothness (FWHM, Friedman 2008), percentage of artifact voxels (QI1, Mortamet 2009), and signal to noise ratio (SNR, Magnotta and Friedman, 2006). Spatial and temporal measures are calculated on the function data. The spatial measures include: EFC, FBER, and FWHM, in addition to ghost to signal ratio (GSR). The temporal measures include: the standardized root mean squared change in fMRI signal between volumes (DVARs; Nichols 2013), mean root mean square deviation (MeanFD, Jenkinson 2003), the percentage of voxels with meanFD ≥ 0.2 (Percent FD; Power 2012), the temporal mean of AFNIs 3dTqual metric (1 minus the Spearman correlation between each fMRI volume and the median volume; Cox 1995) and the average fraction of outliers found in each volume.

2.5 Overall evaluation and the static pricing model

AWS outlines the different service charges in detail; this makes it possible to develop time and cost models for jobs to run on AWS as long as the user is aware of the cluster configuration they'd like to use and the average runtime per job of interest they are submitting. A static pricing model was developed to estimate time and costs for job submissions in Python and Excel. This model breaks the costs down into three categories: compute costs, the costs associated with actual computation and node runtimes, data transfer costs, and storage costs. The model similarly breaks runtimes down into three categories: computing run time, data up transfer time, and data down transfer time. The data transfer and storage costs can vary based on the cluster configuration and storage model being used. For example, data transfer is typically faster when uploading output results to S3 from the cluster, rather than downloading them to a server or workstation. Additionally, storage costs differ between EC2 EBS and S3 storage.

2.6 Impact of spot pricing

To take advantage of the Amazon spot market, another model was developed to pull historical pricing data from AWS and run Monte Carlo simulations to obtain average runtimes and costs for that period of history (03-15 through 09-04 of 2015). The model takes in a bid ratio as an argument and runs Monte Carlo simulations by mocking the pipeline execution incrementally throughout the time period. The bid ratio is expressed as the ratio to the historical mean of the hourly price to run that instance. If the user wants to bid a ratio of 1.5 on an instance with a mean historical spot price of \$0.50, then their bid price will be set to \$0.75. The tuning of this parameter to be a ratio rather than a fixed price helps to control for the mean price fluctuations across time periods and instance types. The model here is similar to the static pricing model except that the hourly price changes with the market as it runs the job submission. In the event that the market price goes above the bid price, the instances have to halt their current processing and resume it when the price dips back below the bid. All of the variables in the Monte Carlo spot pricing model are the same as in the static model except it returns two more quantities of interest: number of interruptions experienced during the run, and the total amount of time spent waiting for the market price to drop down

Table 2. Benchmark results

	Dataset	Dataset size	Platform	Nodes	Parallelization factor	CPU Time	Wall time	Cost (\$)	
Business preprocessing (4 strategies) protocol	ABIDE/NDAR	3197	C-PAC	20	8	23,018	147	760.24	C
	ABIDE	1112	C-PAC	20	3	834	22	80.54	C
	ABIDE	1112	C-PAC	20	4	380	14	19.02	C
	NDAR	986	NITRC-CE	4	32	23,644	193	211.44	C
	NDAR	1247	NITRC-CE	4	32	208	3	2.19	i
	NDAR	1349	NITRC-CE	4	32	450	13	4.69	i

below the bid. With these numbers, the model can provide a probability of forced node termination or the amount of time the user can expect to wait when running their jobs using spot instances.

2.7 The cost of computer ownership

The main arguments for the cloud computing model are having no up-front capital costs and the ability to scale computing resources dynamically. Drawbacks to the model include lack of isolated security oversight of resources on the part of the user and a known, fixed up-front cost for budgeting concerns. It becomes important to breakdown the costs associated with owning ones own servers in order to verify the cloud as a viable solution.

These costs can be broken down to three categories: hardware, maintenance, and energy use. The price for a c3.8xlarge-comparable system (32 cores, 60 GB RAM) is a Dell Precision 7910 with dual Intel Xeon E5-2630 2.4 GHz 8-core processors (32 total virtual CPUs when hyper threading), 64GB of RAM, two 400GB solid state drives (SSD) for local storage, and a 1100 Watt power supply. The costs for this are \$8,642 (from dell.com on 1/31/2015). Maintenance costs include software and hardware maintenance by a research technician; assuming a salary of \$50,000 a year and 5% of their time comes to \$2,500. Energy usage of the server is based on the cost of electricity. The average commercial electricity costs in the United States for November 2014 was \$0.1055 per kilowatt hour. Using this and assuming the server is utilized well at 90% of its power capacity, the annual energy cost comes to \$914.94. The sum total for cost of computer ownership for the first year is \$12,056.94.

These costs dont directly factor in the compute technology becoming obsolete over time, under-utilization, over-utilization (and thus the need to buy more servers), and additional data storage.

3 RESULTS

3.1 Results from benchmarks

The results from the data processing are outlined in the table below. The nodes category indicates the size of the cluster that was run to process the shown dataset. The parallelization factor indicates the number of datasets to be run on a single node at once. CPU time is a measure of the compute cycles needed to process a given dataset, while the wall time is the actual amount of minutes passed during the processing of a given dataset. The cost and cost per dataset come from the total charges incurred on AWS from preprocessing the datasets using the specified pipeline.

These benchmark results yielded realistic estimates for run times, costs, and resources used; these estimates parameterized static and spot-price models that were simulated across a range of bid ratios and dataset sizes.

Table 3. Pipeline runtime parameters

Pipeline	Input size (GB)	Output size (GB)	Parallelization factor	Runtime (mins)	Upload rate (Mb/s)	Download rate (Mb/s)
ANTs	0.007	0.097	4	399	18	20
C-PAC	0.055	2.3	3	33	18	20
Freesurfer	0.007	0.379	8	410.4	18	20

3.2 Results from static simulations

By taking the means of the runtimes, download and upload times, and used storage from the benchmark runs, we were able to parameterize pipeline-level configurations for simulating each pipeline's performance in both a static and spot pricing model. Specifically, we gathered the parameters of the ANTS, C-PAC, and Freesurfer pipelines and ran these through a series of simulations to gauge what scenarios necessitate the use of spot instances versus on-demand in getting the desired results with minimum cost and time.

The static-price model simulated each pipeline's performance across all of the AWS availability zones for a range of dataset sizes. It did this assuming a fixed per-hour runtime cost associated with the desired instance (in our case, the mean of the c3.8xlarge spot price history). This model uses each pipeline's performance configurations - including input and output dataset sizes, parallelization factor, processing time, and upload and download transfer rates - to calculate total runtimes and costs associated with processing data through that pipeline on AWS. The model did the calculations based on:

- 1) uploading the dataset to the master node and storing the data on an EBS volume
- 2) processing the data on as many slave nodes as required for the entire dataset size (up to 20)
- 3) uploading the output data from the pipeline to S3 as the jobs finish
- 4) terminate the cluster when all of the data is finished processing
- 5) store outputs in S3 as long as it takes to download them to a local computer

In order to accurately reflect the costs incurred from processing these datasets, the static model keeps track of the EC2 per-hour runtime costs, data transfer costs, EBS storage costs, and S3 storage costs across all of the availability zones.

3.3 Results from static monte carlo simulations

The spot model was developed with the same pipeline configuration parameters as the static model except that it used the real history of the per-hour spot price of the c3.8xlarge instance (from 3/15/15-9/4/15) during the dataset processing time. In addition to varying the pipeline, dataset size, and availability zone, this model also varied the per-hour bid-price for running the compute nodes. The pipeline would then be run under this bid price; if the spot price for the instance met or exceeded this bid, the processing for the datasets running at the time would be interrupted, and restarted the next time the spot price went below the bid. This simulation was run over and over throughout the history every 20-minutes to yield many results that we can draw statistical conclusions from.

The static and spot models were run across all 23 availability zones on AWS, varying dataset sizes from [50, 100, 200, 500, 1000, 2000, 5000, 10000]. The spot model was run across the bid ratios [1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, 5.0] of the mean per-hour spot history - for the static model the bid ratio was fixed to 1.0.



Figure 1. Enter the caption for your figure here. Repeat as necessary for each of your figures

3.3.1 price and time vs. spot price ratio

The static model has the property that because the compute nodes do not face forced termination, its compute time will always happen as fast or faster than the spot-price model. What is interesting is studying by how much faster it is for the price difference one pays for different sized datasets and for different bid ratios.

It can be seen from the plots that the biggest deciding factor for setting a bid ratio to ensure optimal time/cost for one's goal is dependent on availability zone. However, the cost and runtimes within an availability zone demonstrates a fairly repeatable pattern across pipelines - so it behooves us to focus on the trade-offs of a particular availability zone, or group of availability zones (AWS region).

3.3.2 price and time vs. num datasets

3.3.3 price and time vs. num nodes

3.3.4 price and time vs. instance types

3.4 Monte carlo vs static simulations

4 DISCUSSION

4.1 Comparison of cloud vs. computer ownership

4.2 Strategies for spot bidding

4.3 Finding optimal computing configurations

4.4 Improving performance

5 CONCLUSIONS

Additional Requirements:

DISCLOSURE/CONFLICT-OF-INTEREST STATEMENT

The authors declare that the research was conducted in the absence of any commercial or financial relationships that could be construed as a potential conflict of interest.

The statement about the authors and contributors can be up to several sentences long, describing the tasks of individual authors referred to by their initials and should be included at the end of the manuscript before the References section.

Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text Text
Text Text Text Text Text. Text Text Text Text Text Text Text Text Text Text Text Text Text
Text Text Text Text Text Text Text Text Text Text.

Funding: Text Text Text Text Text Text Text Text.

Supplementary Material should be uploaded separately on submission, if there are Supplementary Figures, please include the caption in the same file as the figure. LaTeX Supplementary Material templates can be found in the Frontiers LaTeX folder

[illegible]

Figure 2. Enter the caption for your figure here. Repeat as necessary for each of your figures