



Harnessing cloud computing for high capacity analysis of neuroimaging data from NDAR

Daniel Clark¹, Christian Haselgrove², David Kennedy², Zhizhong Liu³,
Michael Milham¹, Petros Petrosyan⁴, Carinna Torgerson³, John Van Horn³, Cameron Craddock¹

¹Child Mind Institute, New York, NY, ² University of Massachussetts Medical School, Worcester, MA, ³University of Southern California, Los Angeles, CA, ⁴UCLA, Los Angeles, CA, ⁵Nathan S. Kline Institute for Psychiatric Research, Orangeburg, NY



Introduction

- ▶ The National Database for Autism Research (NDAR)¹ hosts a vast collection of neuroimaging datasets that can be processed and utilized to yield significant scientific discoveries.
- ▶ This amount of resources necessitates a high-performance computing (HPC) infrastructure, which is not always readily available for researchers in-house.
- ▶ Amazon Web Services (AWS) Elastic Compute Cloud (EC2) computing service offers a “pay as you go” model that allows researchers to utilize HPC performance without the up-front capital costs and maintenance of an in-house solution.
- ▶ The developers of the Laboratory of Neuro Imaging (LONI) Pipeline, the Neuroimaging Informatics Tools and Resources Clearinghouse (NITRC) Computational Environment (CE) and the Configurable Pipeline for the Analysis of Connectomes (C-PAC) have implemented pipelines in EC2 that interface with NDAR

Methods

LONI Pipeline

- ▶ The LONI Pipeline software was extended to include new pipeline modules to access data from the NDAR database, transfer input data out of Amazon S3 (Simple Storage Service), and to load results back into S3²
- ▶ A pipeline was constructed to extract cortical thickness and subcortical region volume data from structural MRI images in the NDAR database, which included:
 1. Reorient images to standard orientation using FSL’s reorient2std module
 2. Extract cortical thickness using FreeSurfer recon-all
 3. Calculate volumes of subcortical regions using FSL’s BET and FIRST all
- ▶ The resulting pipeline was used to process 780 T1-weighted structural images and return the results to NDAR

Configurable Pipeline For the Analysis of Connectomes (C-PAC)³

- ▶ C-PAC modules were written in Python to build input data lists by querying NDAR, read input data from S3, write processed results to S3 and write values back to the NDAR database
- ▶ New pipelines were created to perform the ANTS cortical thickness⁴ procedure and the Preprocessed Connectomes Project’s Quality Assessment Protocol (<http://preprocessed-connectomes-project.github.io/quality-assessment-pipeline>)⁵
- ▶ The resulting pipelines and modules were used to process several datasets and return the results to NDAR
 1. Cortical extraction from 3,197 T1-weighted structural images
 2. Structural and functional processing for 1,112 datasets from ABIDE
 3. Automated quality assessment of 1,112 datasets from ABIDE

Neuroimaging Informatics Tools and Resources Clearinghouse Computational Environment (NITRC-CE)⁶

- ▶ The NITRC pipeline processed data using three primary utilities
 1. Extract anatomical and surface-base measures with FreeSurfer recon-all
 2. Segment subcortical structures using FSL’s FIRST to produce volumetric and mesh outputs
 3. Time series QA measures using the fmriqa generate.pl utility from the BXH/XCEDE Tools suite, including mean intensity, center of mass, per-slice variation, and others
- ▶ Python modules were created to query and download data from NDAR as well as to store results back to their database
- ▶ The recon-all and FIRST tools processed 986 and 1,247 T1-weighted anatomical scans, respectively; the fmriqa generate QA measures were generated from 1,349 functional scans.

Interacting with the NDAR database

- ▶ Launched an AWS-hosted miNDAR database by querying NDAR website for the data of interest (e.g. from a particular study)
- ▶ Built a subject list by querying the database for subjects of interest to pass to our pipeline
- ▶ Launch an AWS EC2 HPC cluster using Starcluster⁷
- ▶ Log into the cluster and submit a Sun Grid Engine job using our pipeline software and the subject list
- ▶ The pipeline software will process the data, store raw outputs in an AWS S3 bucket and insert S3 filepaths and output measures into miNDAR database

EC2 Spot Pricing

- ▶ In addition to the “pay as you go” compute pricing model, AWS offers users to bid on instances that are not being used
- ▶ As opposed to paying an “on-demand” fixed price, the hourly charge for the instances will vary over time based on what people in different regions are willing to pay
- ▶ If the spot market price goes above what the user bids, their instance will be terminated by Amazon and all of their processes will be stopped and their data lost
- ▶ Typically, one can expect to pay one-eighth of the hourly price of a standard on-demand instance when using spot-pricing
- ▶ AWS offers cloud solutions based on geographic region, where the price for services vary
- ▶ Some regions fluctuate on spot price more than others, but there are ways to minimize process interruptions and lost data in how one configures their cluster and pipelines

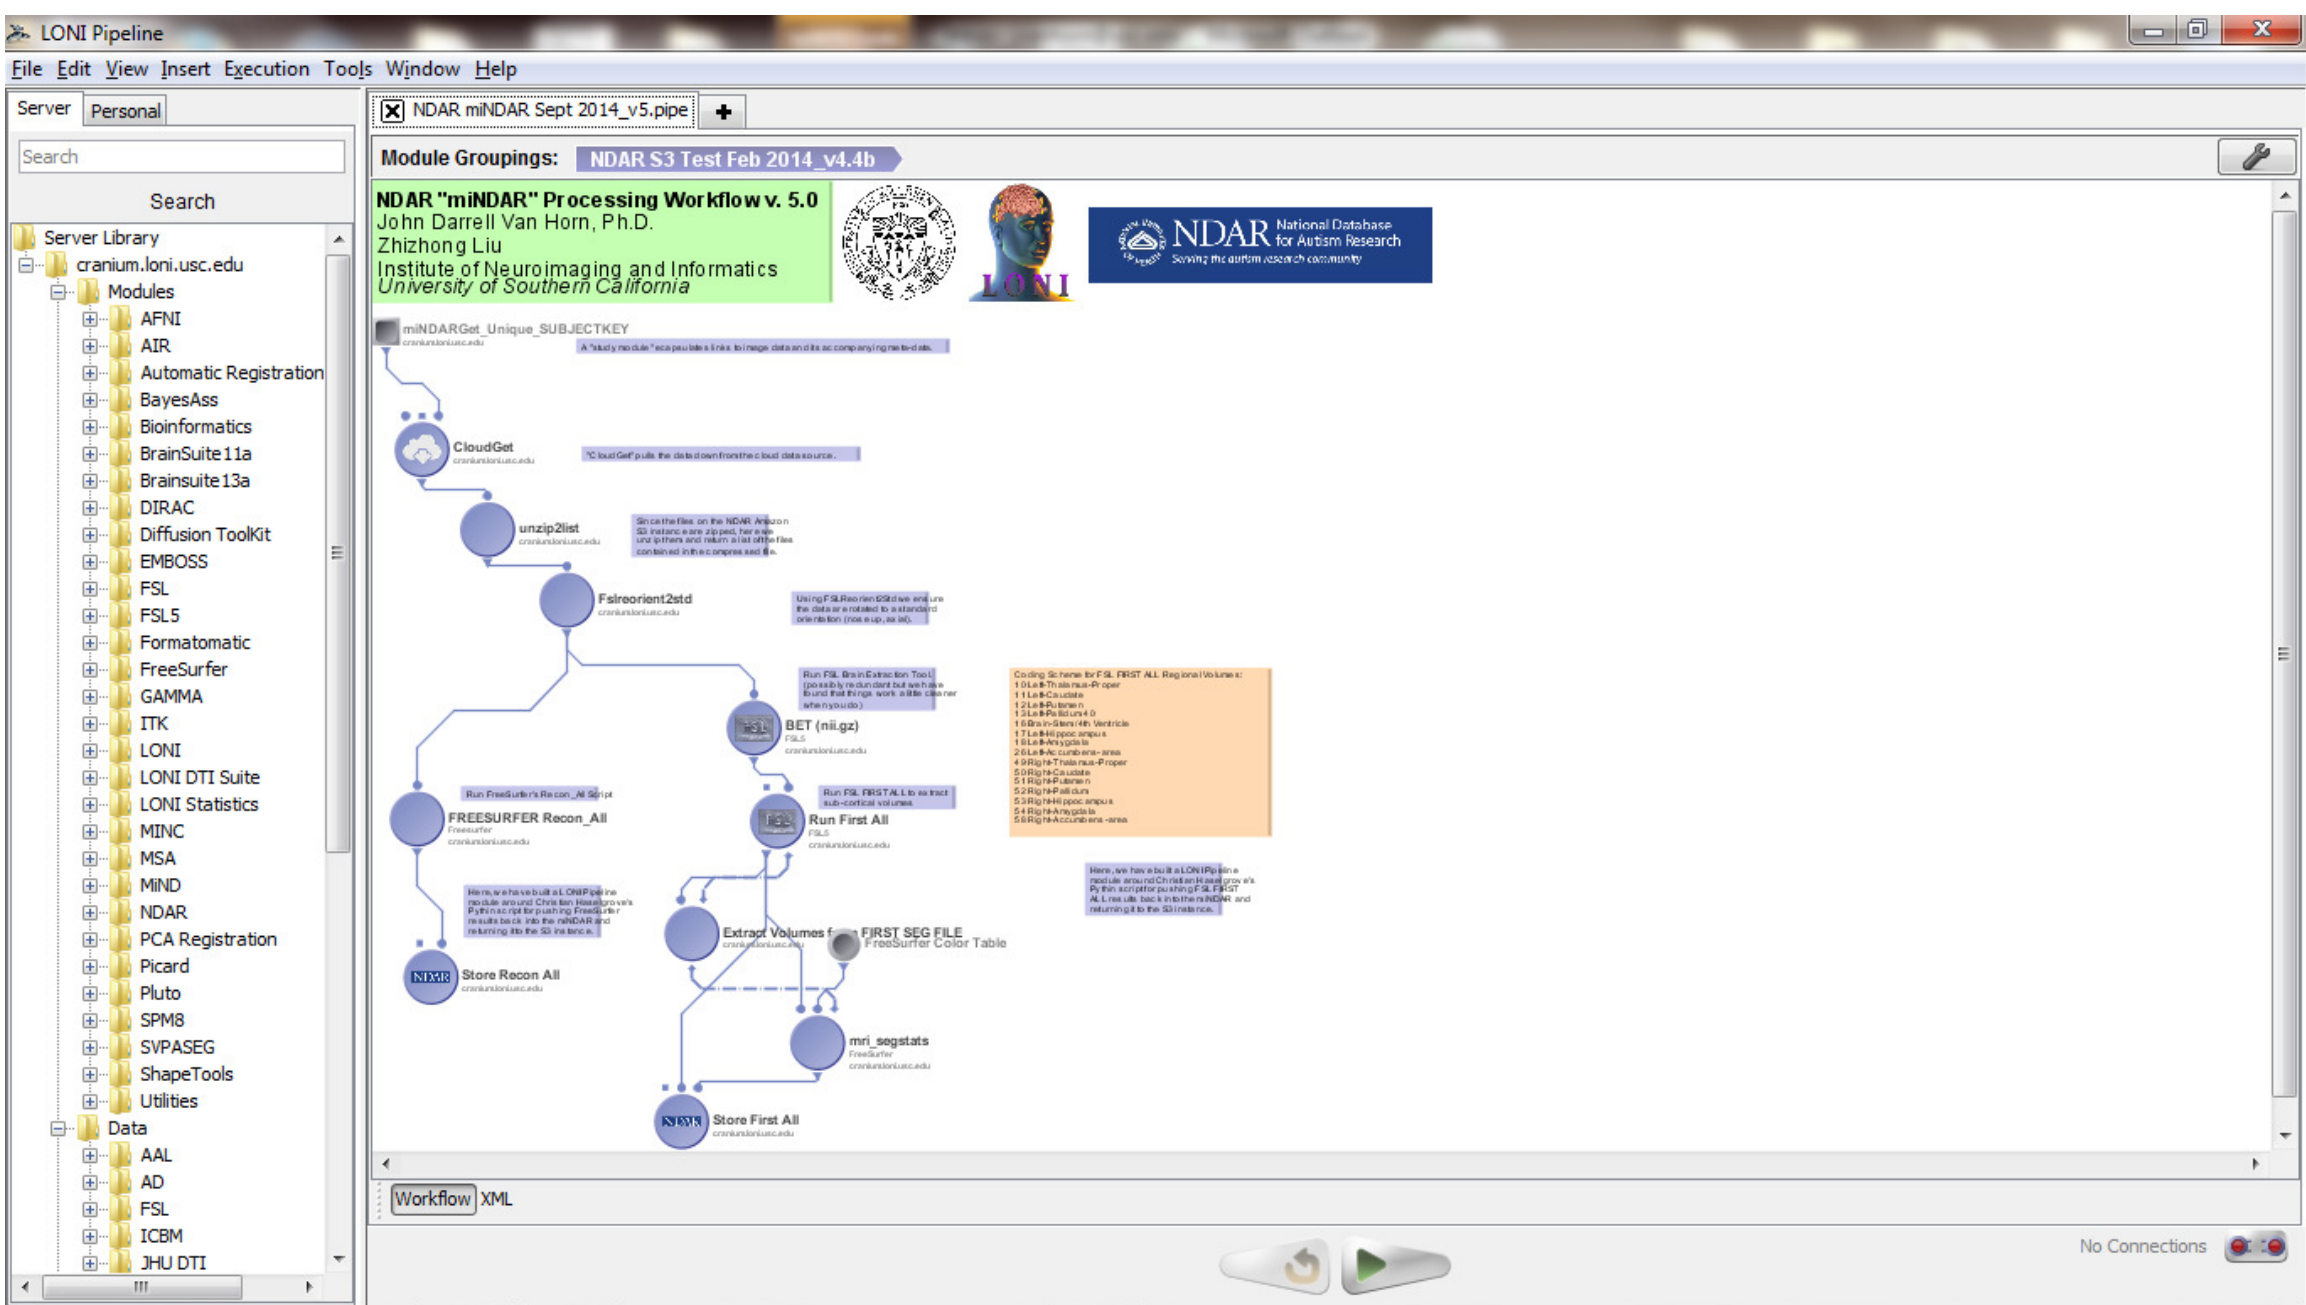


Figure 1 : Graphical layout of the constructed pipeline

Column Name	Data Type	Nullable	Data Default	Column ID	Comments
1 IMAGE_ID	NUMBER(38, 0)	Yes	(null)	1	(null)
2 DATASET_ID	NUMBER(38, 0)	Yes	(null)	2	(null)
3 SUBJECT_ID	VARCHAR2(255 BYTE)	Yes	(null)	3	(null)
4 SPC_SUBJECT_ID	VARCHAR2(255 BYTE)	Yes	(null)	4	(null)
5 INTERVIEW_DATE	TIMESTAMP(6) WITH TIME ZONE	Yes	(null)	5	(null)
6 INTERVIEW_AGE	NUMBER(38, 0)	Yes	(null)	6	(null)
7 GENDER	VARCHAR2(50 BYTE)	Yes	(null)	7	(null)
8 COMMENTS_MISC	VARCHAR2(1000 BYTE)	Yes	(null)	8	(null)
9 IMAGE_FILE	VARCHAR2(255 BYTE)	Yes	(null)	9	(null)
10 IMAGE_THUMBNAIL_FILE	VARCHAR2(255 BYTE)	Yes	(null)	10	(null)
11 IMAGE_DESCRIPTION	VARCHAR2(1512 BYTE)	Yes	(null)	11	(null)
12 IMAGE_FILE_FORMAT	VARCHAR2(50 BYTE)	Yes	(null)	12	(null)
13 IMAGE_MODALITY	VARCHAR2(50 BYTE)	Yes	(null)	13	(null)
14 SCANNER_MANUFACTURER_ID	VARCHAR2(50 BYTE)	Yes	(null)	14	(null)
15 SCANNER_TYPE_ID	VARCHAR2(50 BYTE)	Yes	(null)	15	(null)
16 SCANNER_SOFTWARE_VERSION_ID	VARCHAR2(100 BYTE)	Yes	(null)	16	(null)
17 MAGNETIC_FIELD_STRENGTH	VARCHAR2(50 BYTE)	Yes	(null)	17	(null)
18 HWS_REFERENCE_TIME_ID	PLAT	Yes	(null)	18	(null)
19 HWS_LOCATION_ID	PLAT	Yes	(null)	19	(null)
20 FLIP_ANGLE	VARCHAR2(50 BYTE)	Yes	(null)	20	(null)
21 ACQUISITION_MATRIX	VARCHAR2(50 BYTE)	Yes	(null)	21	(null)

Figure 2 : miNDAR database




Figure 3 : Spot price history for the c3.8xlarge instance in the us-east-1 region on AWS across the four different availability zones for the past three months

Results cont.

Costs and run time models

- ▶ Models assume user uploads their input data to the cloud, runs their pipelines and (optionally) downloads their outputs as they become available.
- ▶ Outputs are stored on an AWS Elastic Block Store (EBS) hard drive mounted to the master node and is NFS-shared across all of the slave nodes.
- ▶ Users can also upload their processed outputs to a cloud storage solution, like AWS Simple Storage Service (S3) directly from the cluster - this is the approach that was taken in processing data for NDAR.
- ▶ Storing results in AWS S3 avoids costly download time and provides for a viable solution for backing up data; however this does incur additional storage costs.

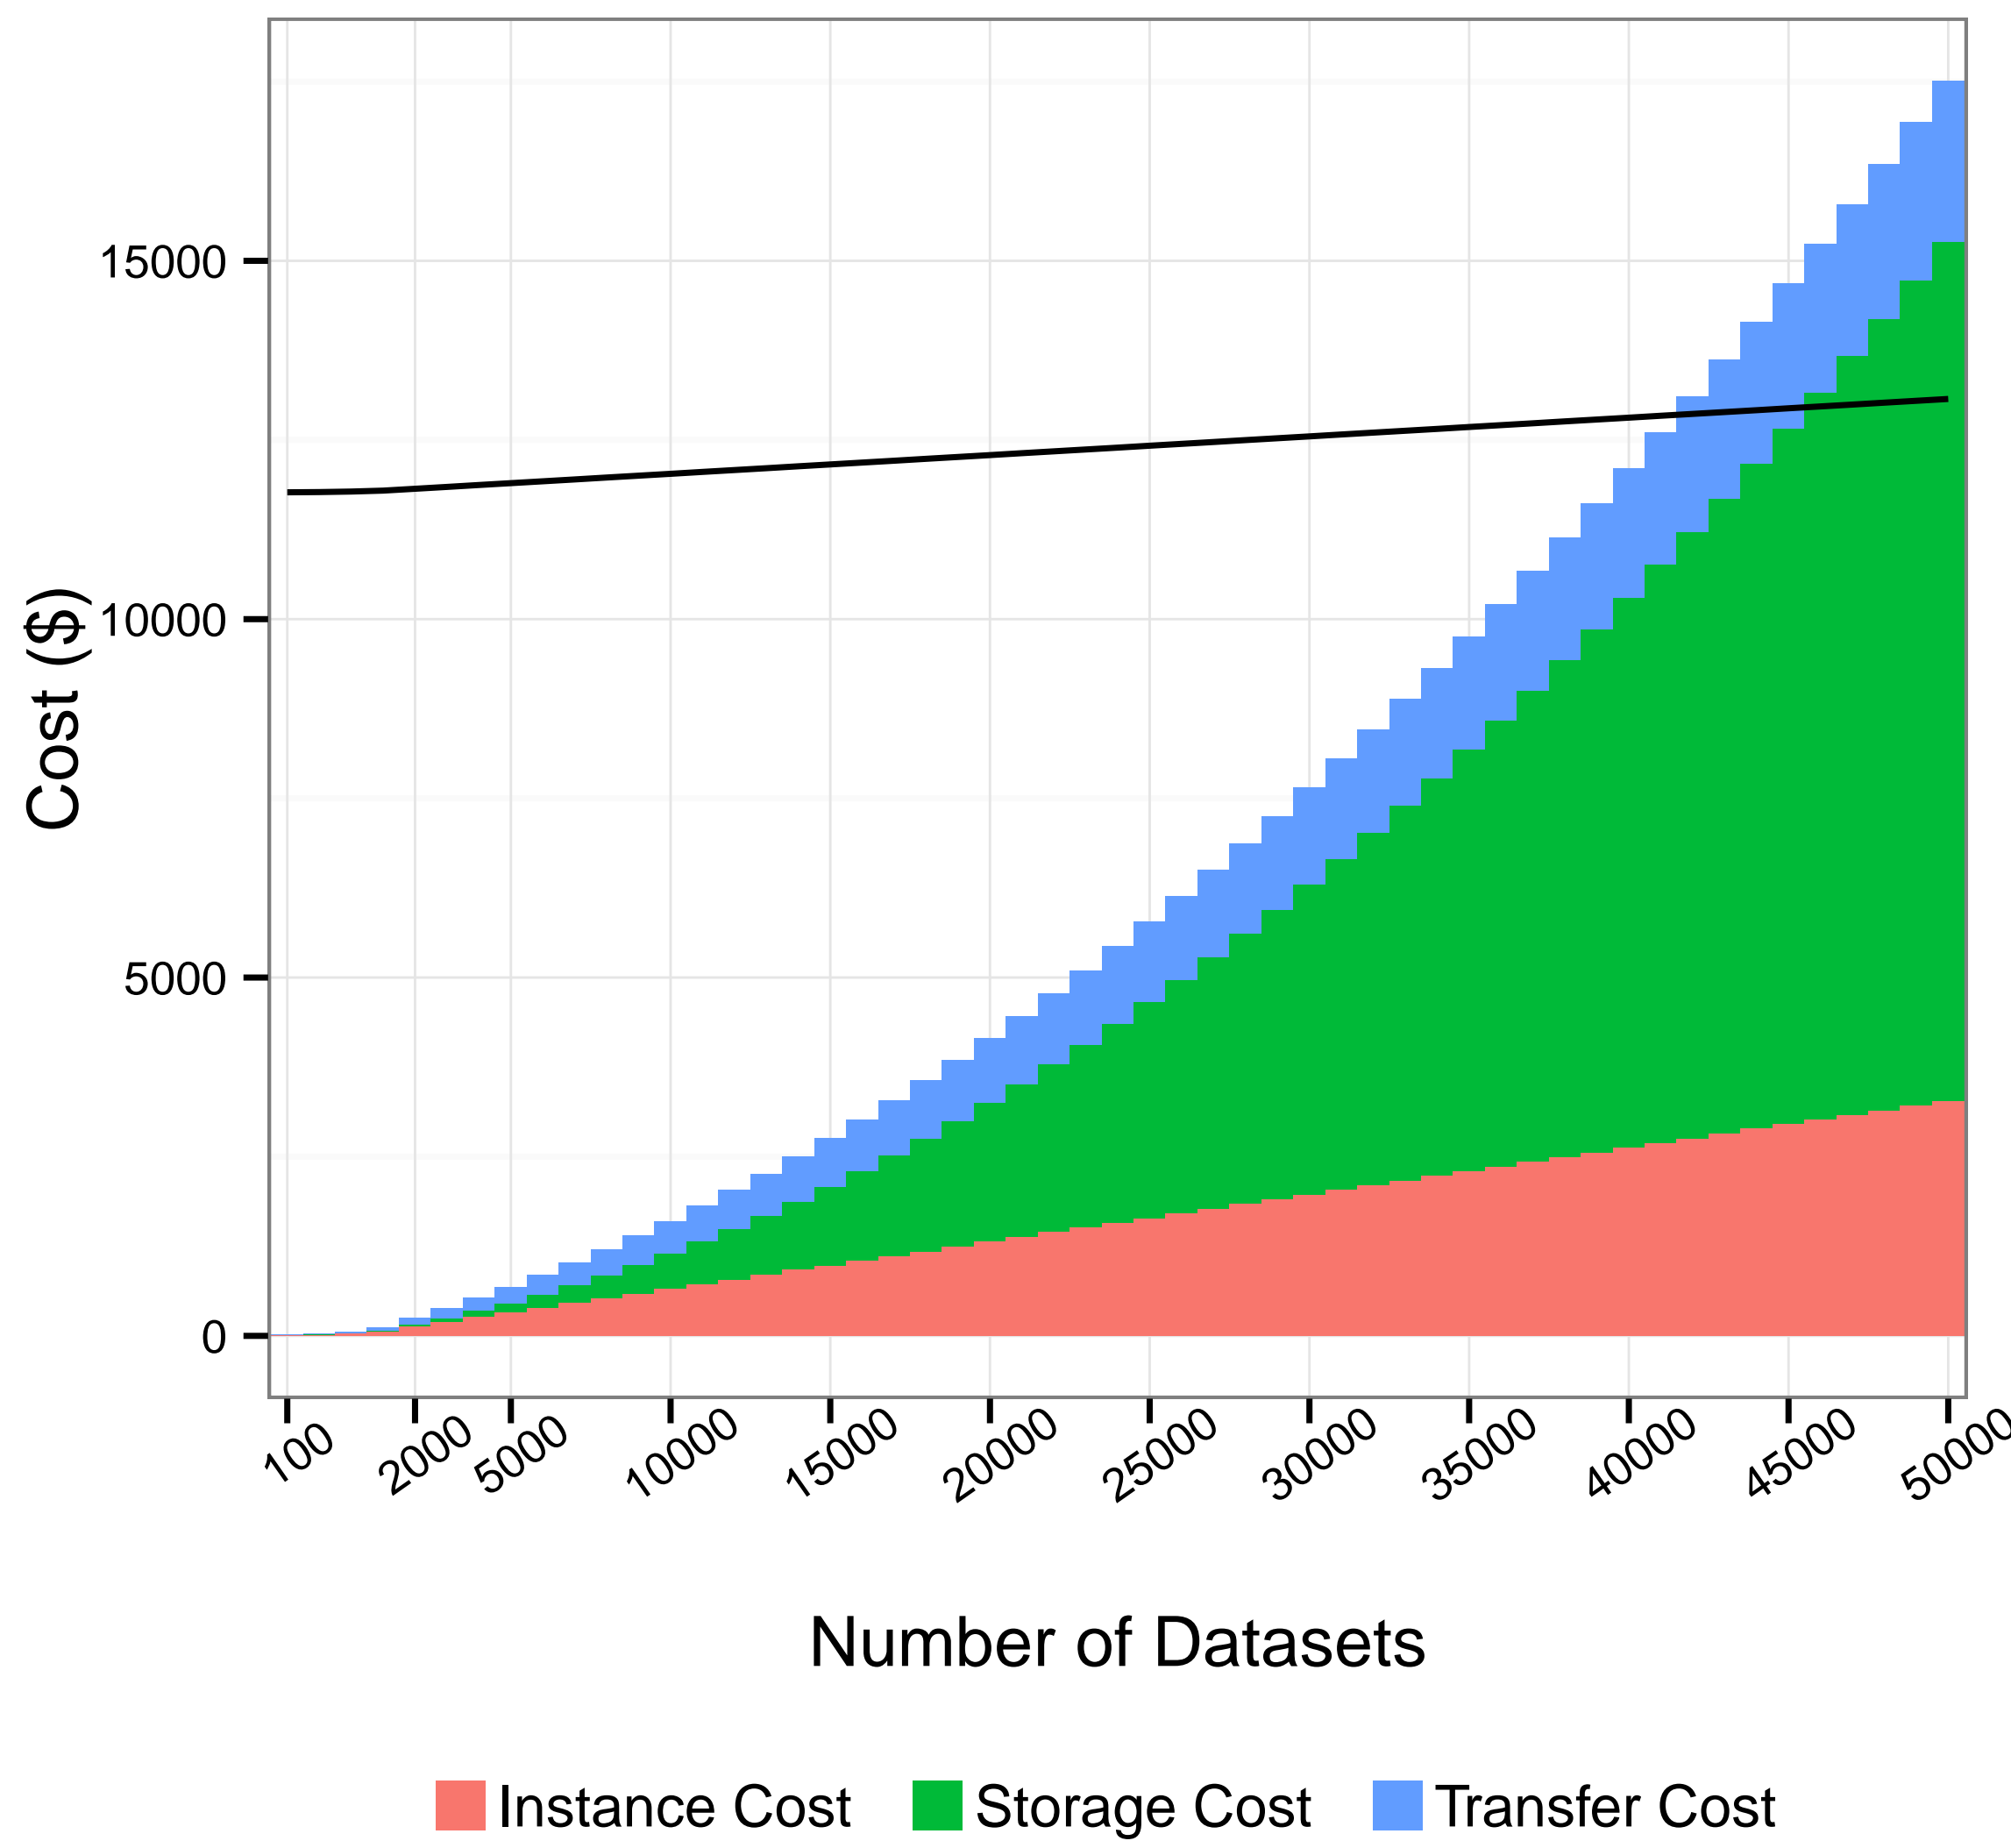


Figure 4 : AWS EC2 costs, grouped by cost type, for a typical C-PAC pipeline for different sized datasets versus owning and maintaining own server; costs are based on using the cluster configuration shown in Table 1 for the cloud, and a single compute node locally




Figure 5 : AWS EC2 costs, grouped by cost type, for a typical FreeSurfer pipeline for different sized datasets versus owning and maintaining own server; costs are based on using the cluster configuration shown in Table 1 for the cloud, and a single compute node locally




Figure 6 : Run times using the C-PAC pipeline, grouped by downloading vs non-downloading output data, for different sized datasets versus running locally; times are based on using the cluster configuration shown in Table 1 for the cloud, and a single compute node locally



Figure 7 : Run times using the FreeSurfer pipeline, grouped by downloading vs non-downloading output data, for different sized datasets versus running locally

Results

Table 1 : Processing completed as a part of the effort. Nodes corresponds to the number of hosts used in the calculation. PF is parallelization factor and corresponds to the number of jobs ran in parallel on each node. On demand instances were used for the master node and spot instances were used for all computation nodes in the cluster. CPU Time is the total amount of time required to perform the computation and Wall Time is the amount of time that passed. # DS: Number of datasets. CPD: Cost Per Dataset. C-PAC: Configurable Pipeline for the Analysis of Connectomes. NITRC-CE: NITRC Computational Environment

Processing	# DS	Platform	Nodes	PF	CPU Time	Wall Time	Cost	CPD
ANTS Cortical Thickness	3197	C-PAC	20	8	23,018	147	\$760.24	\$0.24
Resting state fMRI processing w/ 4 strategies	1112		20	3	834	22	\$80.54	\$0.07
Quality Assessment Protocol	1112		20	4	380	14	\$19.02	\$0.02
FreeSurfer recon-all	986	NITRC-CE	4	32	23,664	193	\$211.44	\$0.21
FSL FIRST	1247		4	32	208	3	\$2.19	> \$0.01
Temporal QA	1349		4	32	450	13	\$4.69	> \$0.01
FreeSurfer recon-all and FSL FIRST	780	LONI Pipeline	20	32	18,720	49	\$252.36	\$0.32

Benefits of the cloud

- ▶ The scalability and payment model for processing this data on a cloud platform allowed for an efficient processing of the NDAR datasets, both in time and cost
- ▶ Additionally, there was no need to install, configure, and maintain machines locally, saving on overhead
- ▶ Competition for computing resources was controlled as every cluster launch was done by the user doing the processing; no fellow researchers needed to login and use additional computing power
- ▶ Cost and time models were produced to demonstrate the benefits of using cloud computing for increasingly large datasets

Conclusion

- ▶ Cloud computing services like AWS EC2 are viable solutions for the processing and analysis of large amounts of data

References and Acknowledgements

1. Hall, D. et al. (2012), Neuroinformatics. 10(4):331-9
2. Torgerson, C.M. et al. (2015), Brain Imaging and Behavior 9:89-103
3. C-PAC ref
4. Das, S.R. et al. (2009), Neuroimage, 45(3): 867-79
5. QAP ref
6. Haselgrove, C. et al. (2014), Front. Neuroinform. 8:52
7. Starcluster ref

Data collection and salary support was provided by a NARSAD Young Investigator Award to RCC.